

Advanced Algorithms

Spring 2023

June 22, 2023 – 14:30–16:30

First Part: Theory Questions

☒ Question 1 (4 points) Consider the following directed, weighted graph, represented by an adjacency matrix where each numerical value represents the weight of the corresponding edge, and where the symbol '-' indicates the absence of the edge between the corresponding vertices.

	s	a	b	c	d
s	-	2	4	-	-
a	-	-	-1	2	-
b	-	-	-	-	4
c	-	-	-	-	2
d	-	-	-	-	-

☒ Draw the graph.

☒ Run the Bellman-Ford algorithm on this graph, using vertex s as the source. You are to return the trace of the execution, i.e. a table with rows indexed by vertices and columns indexed by iteration indexes (starting from 0) where each entry contains the estimated distance between s and that vertex at that iteration.

☒ Question 2 (4 points) For each of the following problems, say whether it is NP-hard or not and, if not, specify the complexity of the best algorithm seen in class.

☒ Maximum independent set

☒ All-pairs shortest paths

☒ Traveling salesperson problem

☒ Graph connectivity

☒ Question 3 (4 points) Define the set cover problem and briefly describe the $O(\log n)$ -approximation algorithm seen in class.

Second Part: Problem Solving

☒ Exercise 1 (9 points) Consider Dijkstra's algorithm seen in class, which returns the lengths of the shortest paths from a source vertex to all other vertices in directed graphs with nonnegative weights:

☒ Explain how to modify Dijkstra's algorithm to return the shortest paths themselves (and not just their lengths).

8✓ Consider the following algorithm for finding shortest paths in a directed graph where edges may have negative weights: add the same large constant to each edge weight so that all the weights become nonnegative, then run Dijkstra's algorithm and return the shortest paths. Is this a valid method? Either prove that it works (i.e., the returned shortest paths are shortest paths in the original graph), or give a counterexample.

8✓ Now let's switch to minimum spanning trees, and do the same: add the same large constant to each edge weight and then run Prim's algorithm. Either prove that the returned solution is a minimum spanning tree of the original graph, or give a counterexample.

8✓ **Exercise 2 (10 points)** Suppose you throw n balls into $\frac{n}{\delta \ln n}$ bins¹ independently and uniformly at random. Applying the following Chernoff bound show that, with high probability, the bin with maximum load (load = number of balls in the bin) contains at most $12 \ln n$ balls. (Hint: focus first on one arbitrary bin and bound the probability of that bin's load exceeding $12 \ln n$...)

Theorem 1. Let X_1, X_2, \dots, X_n be independent indicator random variables such that $E[X_i] = p_i, 0 < p_i < 1$. Let $X = \sum_{i=1}^n X_i$ and $\mu = E[X]$. Then, for $0 < \delta \leq 1$,

$$\Pr(X > (1 + \delta)\mu) \leq e^{-\mu\delta^2/3}.$$

¹Recall that $\ln n = \log_e n$.