

Example of data analysis

Dataset Pollution

Data Mining
Master Degree in Computer Science
University of Padova

a.y. 2022/2023

Annamaria Guolo

These notes are an example of data analysis using the techniques illustrated during the course. This is just an example: as you know, there is no a unique way to analyse the data and the right model does not exist. Whichever the direction you take for the analysis, please be sure to apply correctly the appropriate instruments for the available data.

1 Dataset Pollution

Dataset included in `pollution.RData` refers to a study about the relationship between mortality in 60 US areas and air pollution. Some environmental and demographical information are collected.

- `mortality`: mortality rate (annual deaths for 100000 persons)
- `precipitation`: mean annual precipitation (inches)
- `humidity`: percent relative humidity
- `Jan.temp`: mean January temperature (Fahrenheit)
- `July.temp`: mean July temperature (Fahrenheit)
- `over65`: percentage of the population aged 65 years or over
- `house`: population per household
- `education`: median number of school years completed for persons 25 years or older

- comfort: percentage of the housing that is sound with all facilities
- density: population density (in persons per square mile)
- office: percentage of office workers
- poor: percentage of households with annual income under 3000 dollars
- HC: level of hydrocarbons
- NOX: dangerous level of oxides of nitrogen?: Si (Yes) ($> 30 \mu g/mc$), No ($\geq 30 \mu g/mc$)
- SO2: dangerous level of sulfur dioxide?: Si (Yes) ($> 125 \mu g/mc$), No ($\geq 125 \mu g/mc$)

FIRST QUESTION.

Consider the dataset composed by mortality, precipitation, humidity, HC, NOX, SO2. Construct the most appropriate model for the purpose of the analysis. Which variables are associated to the mortality rate?

The response variable is mortality rate, so consider a linear regression model.

```
load('pollution.RData')
dim(pollution)

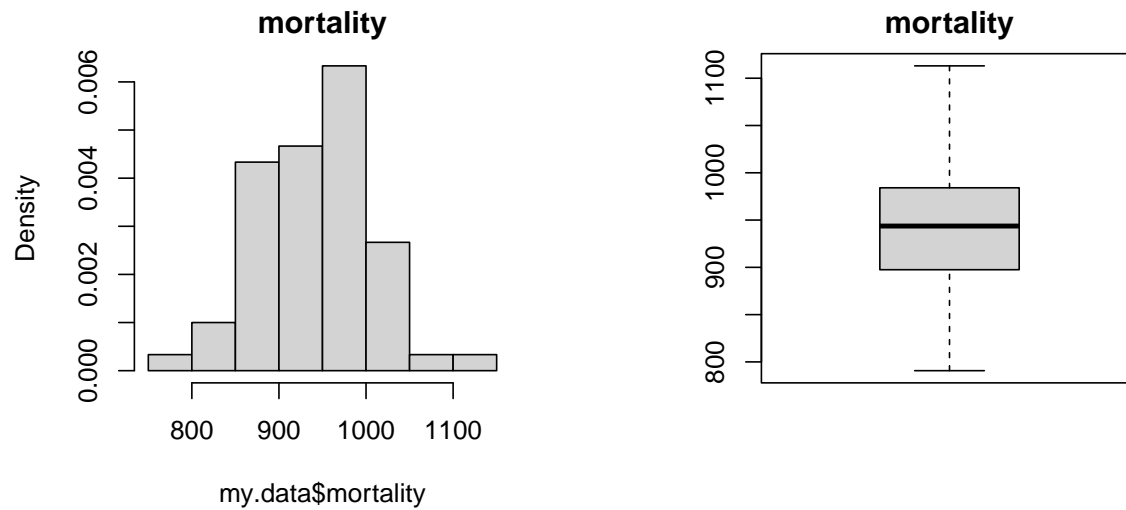
## [1] 60 15

my.data <- pollution[,c('mortality', 'HC', 'NOX', 'SO2', 'precipitation',
                        'humidity')]
dim(my.data)

## [1] 60 6
```

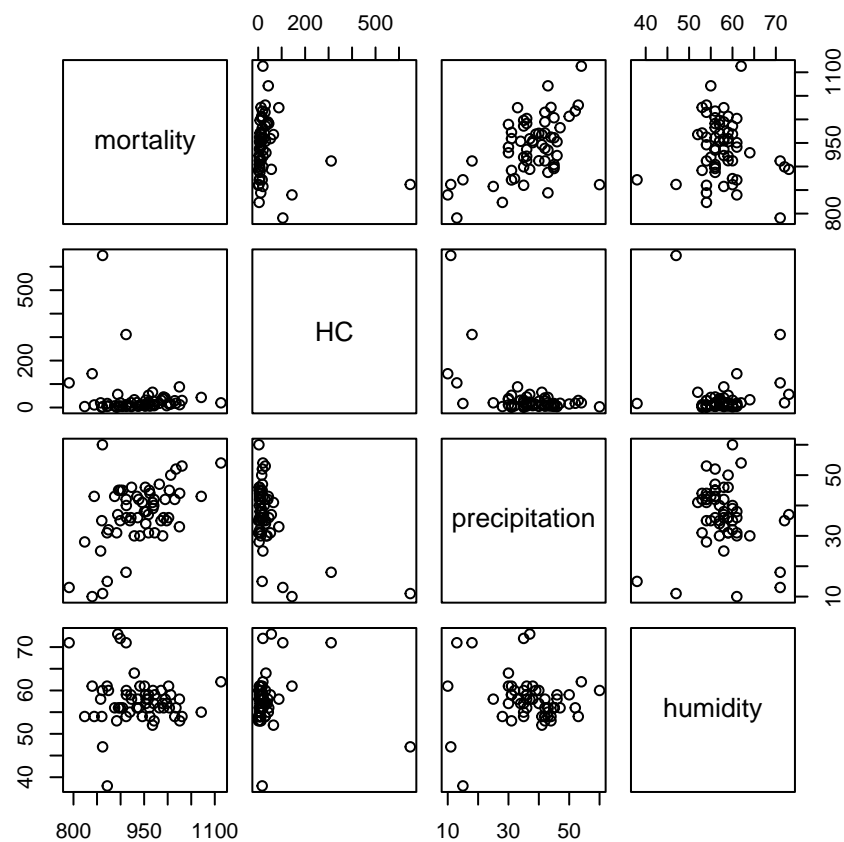
Start with some graphical analyses, working on the response variable, to check the normality of its distribution

```
par(mfrow=c(1,2))
hist(my.data$mortality, prob=TRUE, main='mortality')
boxplot(my.data$mortality, main='mortality')
```




The hypothesis of normality seems to be satisfied.
Consider some graphs to evaluate relationships between variables.

```
pairs(my.data[,c(1,2,5,6)])
```

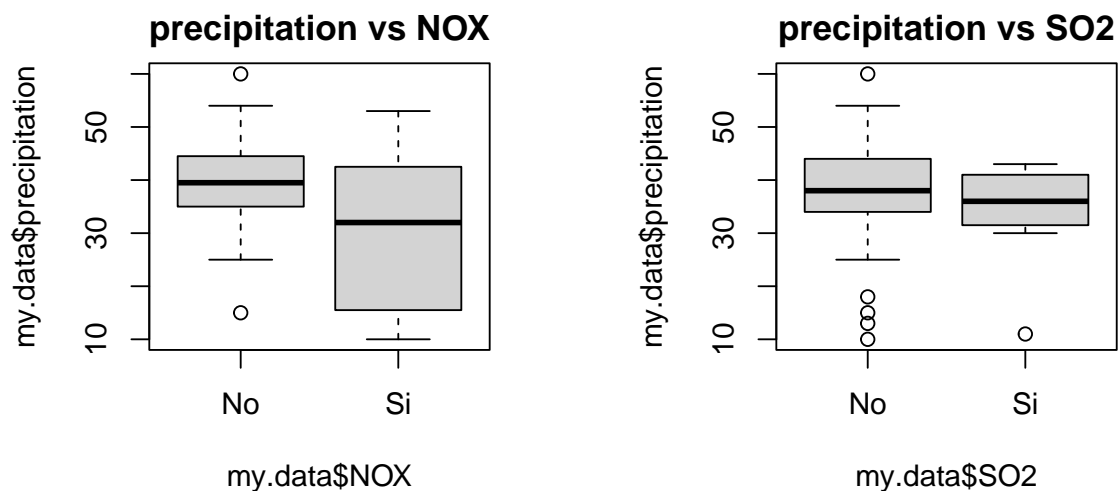


Graphs of interaction between factors and quantitative variables

```
par(mfrow=c(1,2))
boxplot(my.data$mortality~my.data$NOX, main='mortality vs NOX')
boxplot(my.data$mortality~my.data$SO2, main='mortality vs SO2')
```



```
par(mfrow=c(1,2))
boxplot(my.data$precipitation~my.data$NOX, main='precipitation vs NOX')
boxplot(my.data$precipitation~my.data$SO2, main='precipitation vs SO2')
```





```
par(mfrow=c(1,2))
plot(my.data$precipitation, my.data$mortality, col=my.data$NOX,
     main='mortality vs precipitation', ylab='mortality',
     xlab='precipitation')
legend('topleft', legend=c('NOX = No', 'NOX = Yes'), col=c(1,2), pch=c(19,19),
```

```

cex=0.7, bty='n')
plot(my.data$precipitation, my.data$mortality, col=my.data$SO2,
     main='mortality vs precipitation', ylab='mortality',
     xlab='precipitation')
legend('topleft', legend=c('SO2 = No', 'SO2 = Yes'), col=c(1,2), pch=c(19,19),
     cex=0.7, bty='n')

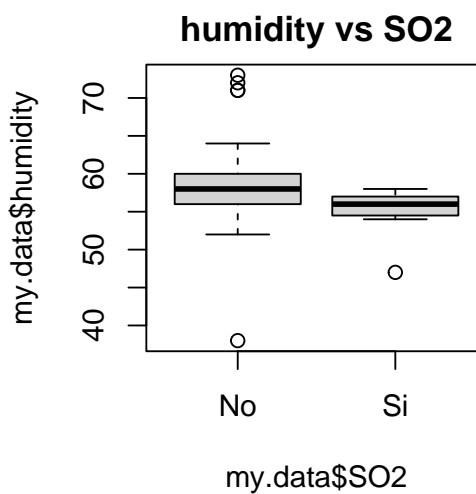

```



```

par(mfrow=c(1,2))
boxplot(my.data$humidity~my.data$NOX, main='humidity vs NOX')
boxplot(my.data$humidity~my.data$SO2, main='humidity vs SO2')

```



```



par(mfrow=c(1,2))
plot(my.data$humidity, my.data$mortality, col=my.data$NOX, main='mortality vs humidity',
     ylab='mortality', xlab='humidity')

```

```

legend('topleft', legend=c('NOX = No', 'NOX = Yes'), col=c(1,2), pch=c(19,19),
      cex=0.7, bty='n')
plot(my.data$humidity, my.data$mortality, col=my.data$SO2, main='mortality vs humidity',
      ylab='mortality', xlab='humidity')
legend('topleft', legend=c('SO2 = No', 'SO2 = Yes'), col=c(1,2), pch=c(19,19),
      cex=0.7, bty='n')



```



```

par(mfrow=c(1,2))
boxplot(my.data$HC~my.data$NOX, main='HC vs NOX')
boxplot(my.data$HC~my.data$SO2, main='HC vs SO2')

```



```

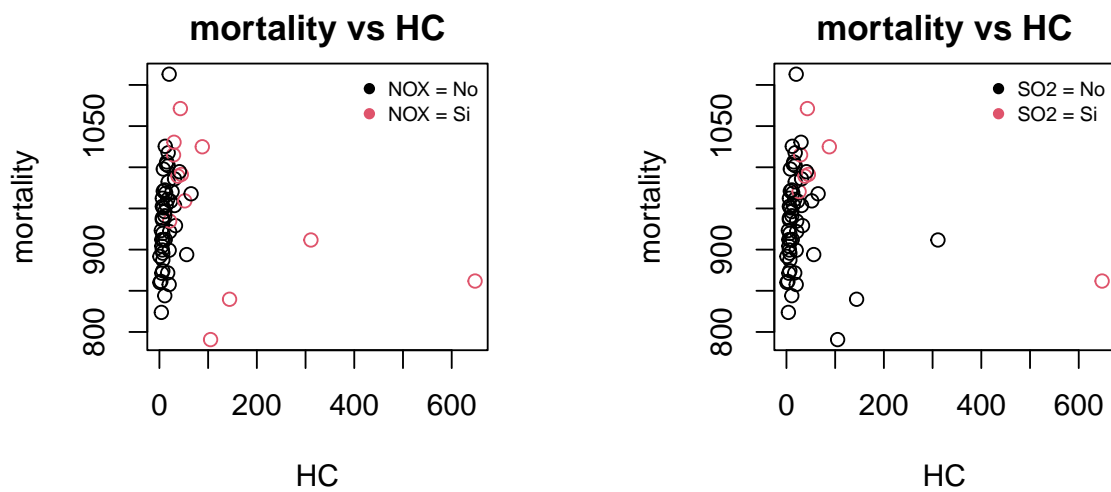
par(mfrow=c(1,2))
plot(my.data$HC, my.data$mortality, col=my.data$NOX, main='mortality vs HC',
      ylab='mortality', xlab='HC')

```

```

legend('topright', legend=c('NOX = No', 'NOX = Si'), col=c(1,2), pch=c(19,19),
      cex=0.7, bty='n')
plot(my.data$HC, my.data$mortality, col=my.data$SO2, main='mortality vs HC',
      ylab='mortality', xlab='HC')
legend('topright', legend=c('SO2 = No', 'SO2 = Si'), col=c(1,2), pch=c(19,19),
      cex=0.7, bty='n')

```




Mosaicplot for factors

```


mosaicplot(table(my.data$NOX, my.data$SO2), las=1,
            cex.axis=1, main='NOX vs SO2')

```



Boxplot of mortality with respect to factors

```
boxplot(my.data$mortality ~ my.data$NOX*my.data$SO2, ylab='mortality', xlab='NOX:SO2')
```



it seems some interesting relationships can be inserted in the model.
Start with a model with interactions between covariates

```
m <- lm(mortality ~ HC*NOX + HC*SO2+ NOX*SO2 +
        precipitation*NOX + humidity*NOX + precipitation*SO2 +
        humidity*SO2, data=my.data)
summary(m)
```

```
##
## Call:
## lm(formula = mortality ~ HC * NOX + HC * SO2 + NOX * SO2 + precipitation *
##      NOX + humidity * NOX + precipitation * SO2 + humidity * SO2,
##      data = my.data)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -123.70  -23.30    0.94   21.46  125.23
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   829.8693    85.9000   9.661 9.64e-13 ***
## HC              1.3385     0.5693   2.351  0.0229 *
## NOXSi          213.5220   389.9062   0.548  0.5865
## SO2Si        -1269.4260  1324.4620  -0.958  0.3427
## precipitation    3.1382     0.9241   3.396  0.0014 **
## humidity       -0.6172     1.4865  -0.415  0.6799
## HC:NOXSi       -0.8982     0.6493  -1.383  0.1731
## HC:SO2Si       -0.1257     0.6046  -0.208  0.8362
```



```
## NOXSi:SO2Si      84.8317    65.7880    1.289    0.2035
## NOXSi:precipitation  1.3275    2.2492    0.590    0.5579
## NOXSi:humidity    -4.3185    5.9624   -0.724    0.4725
## SO2Si:precipitation  3.6625    6.8127    0.538    0.5934
## SO2Si:humidity    20.0289    19.4246    1.031    0.3078
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 48.62 on 47 degrees of freedom
## Multiple R-squared:  0.5132, Adjusted R-squared:  0.389
## F-statistic:  4.13 on 12 and 47 DF,  p-value: 0.0002117
```

Variable selection

```
m2 <- lm(mortality ~ HC*NOX + NOX*SO2 +precipitation*NOX +
          humidity*NOX + humidity*SO2, data=my.data)
summary(m2)

##
## Call:
## lm(formula = mortality ~ HC * NOX + NOX * SO2 + precipitation *
##      NOX + humidity * NOX + humidity * SO2, data = my.data)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -123.70  -25.45    1.10   23.91  125.23
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    829.8693    85.1871   9.742 4.76e-13 ***
## HC              1.3385     0.5645   2.371  0.02171 *
## NOXSi          -43.8206   305.5974  -0.143  0.88657
## SO2Si          -912.1087   757.5886  -1.204  0.23439
## precipitation    3.1382     0.9164   3.424  0.00125 **
## humidity        -0.6172     1.4742  -0.419  0.67727
## HC:NOXSi        -1.1201     0.5988  -1.871  0.06738 .
## NOXSi:SO2Si      68.1618    60.6394   1.124  0.26647
## NOXSi:precipitation  2.2077     2.0670   1.068  0.29072
## NOXSi:humidity   -0.1698     4.4568  -0.038  0.96977
## SO2Si:humidity   16.2730    13.2630   1.227  0.22570
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 48.22 on 49 degrees of freedom
## Multiple R-squared:  0.5009, Adjusted R-squared:  0.3991
```

```
## F-statistic: 4.918 on 10 and 49 DF, p-value: 6.578e-05

m3 <- lm(mortality ~ HC*NOX + NOX*S02+precipitation*NOX +
          humidity*S02, data=my.data)
summary(m3)

##
## Call:
## lm(formula = mortality ~ HC * NOX + NOX * S02 + precipitation *
##      NOX + humidity * S02, data = my.data)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -123.674  -25.389    1.129   23.937  125.271
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)      830.8328    80.5286  10.317 5.56e-14 ***
## HC                1.3405     0.5566   2.408  0.01976 *
## NOXSi            -55.0829    76.5293  -0.720  0.47502
## S02Si            -896.2700   626.9337  -1.430  0.15905
## precipitation     3.1401     0.9057   3.467  0.00109 **
## humidity         -0.6358     1.3772  -0.462  0.64634
## HC:NOXSi         -1.1229     0.5883  -1.909  0.06203 .
## NOXSi:S02Si       68.5023    59.3748   1.154  0.25410
## NOXSi:precipitation 2.2422     1.8381   1.220  0.22825
## S02Si:humidity    15.9946    10.9556   1.460  0.15056
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 47.73 on 50 degrees of freedom
## Multiple R-squared:  0.5009, Adjusted R-squared:  0.4111
## F-statistic: 5.576 on 9 and 50 DF, p-value: 2.633e-05

m4 <- lm(mortality ~ HC*NOX + precipitation*NOX +
          humidity*S02, data=my.data)
summary(m4)

##
## Call:
## lm(formula = mortality ~ HC * NOX + precipitation * NOX + humidity *
##      S02, data = my.data)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
```

```
## -123.004 -27.043 4.352 24.727 126.907
##
## Coefficients:
## Estimate Std. Error t value Pr(>|t|)
## (Intercept) 832.1736 80.7811 10.302 4.64e-14 ***
## HC 1.2771 0.5557 2.298 0.02569 *
## NOXSi -45.0271 76.2778 -0.590 0.55759
## precipitation 3.1187 0.9085 3.433 0.00119 **
## humidity -0.6447 1.3817 -0.467 0.64275
## SO2Si -643.8957 589.4376 -1.092 0.27980
## HC:NOXSi -1.0976 0.5898 -1.861 0.06851 .
## NOXSi:precipitation 2.2898 1.8436 1.242 0.21991
## humidity:SO2Si 12.3731 10.5303 1.175 0.24545
## ---
## Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 47.89 on 51 degrees of freedom
## Multiple R-squared: 0.4876, Adjusted R-squared: 0.4072
## F-statistic: 6.067 on 8 and 51 DF, p-value: 1.781e-05

m5 <- lm(mortality ~ HC*NOX + precipitation*NOX + humidity +SO2, data=my.data)
summary(m5)

##
## Call:
## lm(formula = mortality ~ HC * NOX + precipitation * NOX + humidity +
## SO2, data = my.data)
##
## Residuals:
## Min 1Q Median 3Q Max
## -123.784 -27.825 3.738 29.109 125.486
##
## Coefficients:
## Estimate Std. Error t value Pr(>|t|)
## (Intercept) 811.4126 79.1131 10.256 4.31e-14 ***
## HC 1.2560 0.5575 2.253 0.02850 *
## NOXSi -11.4834 70.9925 -0.162 0.87213
## precipitation 3.0837 0.9113 3.384 0.00137 **
## humidity -0.2497 1.3450 -0.186 0.85345
## SO2Si 48.0313 25.8178 1.860 0.06849 .
## HC:NOXSi -1.2627 0.5749 -2.196 0.03254 *
## NOXSi:precipitation 1.6480 1.7673 0.932 0.35539
## ---
## Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
```

```
## Residual standard error: 48.06 on 52 degrees of freedom
## Multiple R-squared:  0.4738, Adjusted R-squared:  0.4029
## F-statistic: 6.687 on 7 and 52 DF,  p-value: 1.178e-05

m6 <- lm(mortality ~ HC*NOX + precipitation + humidity +SO2, data=my.data)
summary(m6)

##
## Call:
## lm(formula = mortality ~ HC * NOX + precipitation + humidity +
##      SO2, data = my.data)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -131.733  -25.461    4.649   29.974  120.225
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   817.8220    78.7168  10.389 2.17e-14 ***
## HC              1.3178     0.5528   2.384  0.0207 *
## NOXSi          50.6577    24.4451   2.072  0.0431 *
## precipitation  3.5568     0.7561   4.704 1.86e-05 ***
## humidity      -0.7002     1.2537  -0.558  0.5789
## SO2Si          50.9589    25.5946   1.991  0.0516 .
## HC:NOXSi      -1.3881     0.5582  -2.487  0.0161 *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 48.01 on 53 degrees of freedom
## Multiple R-squared:  0.465, Adjusted R-squared:  0.4044
## F-statistic: 7.676 on 6 and 53 DF,  p-value: 5.843e-06

m7 <- lm(mortality ~ HC*NOX + precipitation+SO2, data=my.data)
summary(m7)

##
## Call:
## lm(formula = mortality ~ HC * NOX + precipitation + SO2, data = my.data)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -134.715  -27.120    3.017   28.223  117.194
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
```

```
## (Intercept) 777.6338 31.6991 24.532 < 2e-16 ***
## HC          1.2464 0.5344 2.332 0.0234 *
## NOXSi       46.8190 23.3089 2.009 0.0496 *
## precipitation 3.5797 0.7502 4.772 1.43e-05 ***
## SO2Si       55.0097 24.3884 2.256 0.0282 *
## HC:NOXSi    -1.3126 0.5382 -2.439 0.0180 *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 47.7 on 54 degrees of freedom
## Multiple R-squared:  0.4618, Adjusted R-squared:  0.412
## F-statistic: 9.267 on 5 and 54 DF, p-value: 2.05e-06
```

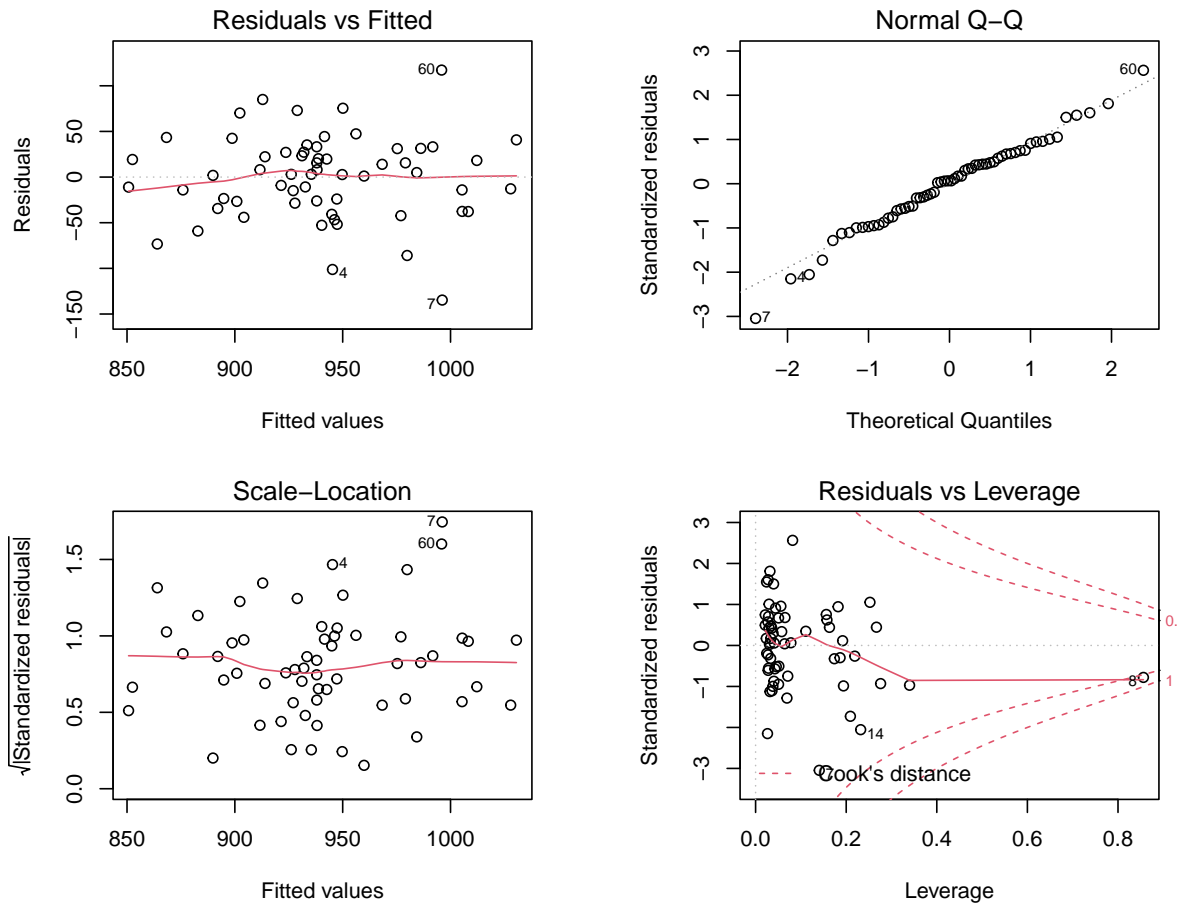
Check the real advantage of using model m7

```
anova(m7, m)

## Analysis of Variance Table
##
## Model 1: mortality ~ HC * NOX + precipitation + SO2
## Model 2: mortality ~ HC * NOX + HC * SO2 + NOX * SO2 + precipitation *
##          NOX + humidity * NOX + precipitation * SO2 + humidity * SO2
##   Res.Df    RSS Df Sum of Sq    F Pr(>F)
## 1      54 122857
## 2      47 111114  7      11744 0.7096 0.664
```

Residual analysis of model m7

```
par(mfrow=c(2,2))
plot(m7)
```



The graphical analysis seems satisfactory: there are no anomalous values, neither trend in residuals or deviations from normality. We can check the need for a polynomial associated to precipitation

```
m8 <- update(m7, .~. + I(precipitation^2))
summary(m8)
```

```
##
## Call:
## lm(formula = mortality ~ HC + NOX + precipitation + SO2 + I(precipitation^2) +
##     HC:NOX, data = my.data)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -103.486  -29.504    3.221   27.981  134.836
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   685.69395    69.28985    9.896 1.21e-13 ***
## HC              1.19891     0.52944    2.265  0.0277 *
```

```
## NOXSi          54.73688    23.65753    2.314    0.0246 *
## precipitation    8.94696     3.68284    2.429    0.0185 *
## SO2Si          41.24401    25.83250    1.597    0.1163
## I(precipitation^2) -0.07359    0.04946   -1.488    0.1427
## HC:NOXSi       -1.19035     0.53853   -2.210    0.0314 *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 47.17 on 53 degrees of freedom
## Multiple R-squared:  0.4834, Adjusted R-squared:  0.4249
## F-statistic: 8.265 on 6 and 53 DF,  p-value: 2.48e-06

## it seems the polynomial is not useful
```

Model m7 suggests that

- the mortality rate increases with precipitation
- the mortality rate increases with SO2
- for a small level of NOX, the mortality rate increases with HC, and viceversa.

Try to see if the model can be improved using splines.

```
library(splines)
sp.HC <- smooth.spline(x=my.data$HC, y=my.data$mortality, cv=TRUE)

## Warning in smooth.spline(x = my.data$HC, y = my.data$mortality, cv = TRUE):
## cross-validation with non-unique 'x' values seems doubtful

## df=33
sp.precipitation <- smooth.spline(x=my.data$precipitation,
                                  y=my.data$mortality, cv=TRUE)

## Warning in smooth.spline(x = my.data$precipitation, y = my.data$mortality,
## : cross-validation with non-unique 'x' values seems doubtful

## df=2
library(gam)

## Loading required package: foreach
## Loaded gam 1.20

m.gam <- gam(mortality ~ s(HC, 33)*NOX + SO2 +
             s(precipitation,2), data=my.data)
summary(m.gam)
```

```
##
## Call: gam(formula = mortality ~ s(HC, 33) * NOX + SO2 + s(precipitation,
##      2), data = my.data)
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -1.049e+02 -7.880e+00  2.680e-07  5.141e+00  9.413e+01
##
## (Dispersion Parameter for gaussian family taken to be 2371.666)
##
##      Null Deviance: 228275.4 on 59 degrees of freedom
## Residual Deviance: 49804.92 on 21 degrees of freedom
## AIC: 653.5642
##
## Number of Local Scoring Iterations: NA
##
## Anova for Parametric Effects
##              Df Sum Sq Mean Sq F value    Pr(>F)
## s(HC, 33)      1   8253     8253  3.4799 0.0761510 .
## NOX            1   1722     1722  0.7261 0.4037529
## SO2            1  36657    36657 15.4560 0.0007653 ***
## s(precipitation, 2) 1  51391    51391 21.6689 0.0001359 ***
## s(HC, 33):NOX    1  15286    15286  6.4452 0.0191057 *
## Residuals      21  49805     2372
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Anova for Nonparametric Effects
##              Npar Df  Npar F Pr(F)
## (Intercept)
## s(HC, 33)      32 0.98998 0.521
## NOX
## SO2
## s(precipitation, 2) 1 0.95727 0.339
## s(HC, 33):NOX
```

We do not need splines for HC.

```
m2.gam <- gam(mortality ~ HC*NOX + SO2 + s(precipitation,2), data=my.data)
summary(m2.gam)

##
## Call: gam(formula = mortality ~ HC * NOX + SO2 + s(precipitation, 2),
##      data = my.data)
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
```



```
## -118.557 -29.873 2.896 30.243 124.691
##
## (Dispersion Parameter for gaussian family taken to be 2237.168)
##
## Null Deviance: 228275.4 on 59 degrees of freedom
## Residual Deviance: 118569.8 on 53 degrees of freedom
## AIC: 641.6074
##
## Number of Local Scoring Iterations: NA
##
## Anova for Parametric Effects
##
```

	Df	Sum Sq	Mean Sq	F value	Pr(>F)
HC	1	4424	4424	1.9776	0.16548
NOX	1	10645	10645	4.7581	0.03362 *
SO2	1	19817	19817	8.8579	0.00439 **
s(precipitation, 2)	1	58131	58131	25.9841	4.708e-06 ***
HC:NOX	1	12322	12322	5.5078	0.02270 *
Residuals	53	118570	2237		

```
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Anova for Nonparametric Effects
##
```

	Npar	Df	Npar F	Pr(F)
(Intercept)				
HC				
NOX				
SO2				
s(precipitation, 2)	1	1.9165	0.172	
HC:NOX				

We do not need splines for precipitation.


SECOND QUESTION.

Consider all the variables in the dataset. Construct the most appropriate model for the purpose of the analysis. Which variables are associated to the mortality rate?

Given the large number of covariates, we can use regularization techniques. And we can exploits some of the previous findings, as, for example, we can insert in the analysis the interaction between HC and NOX.

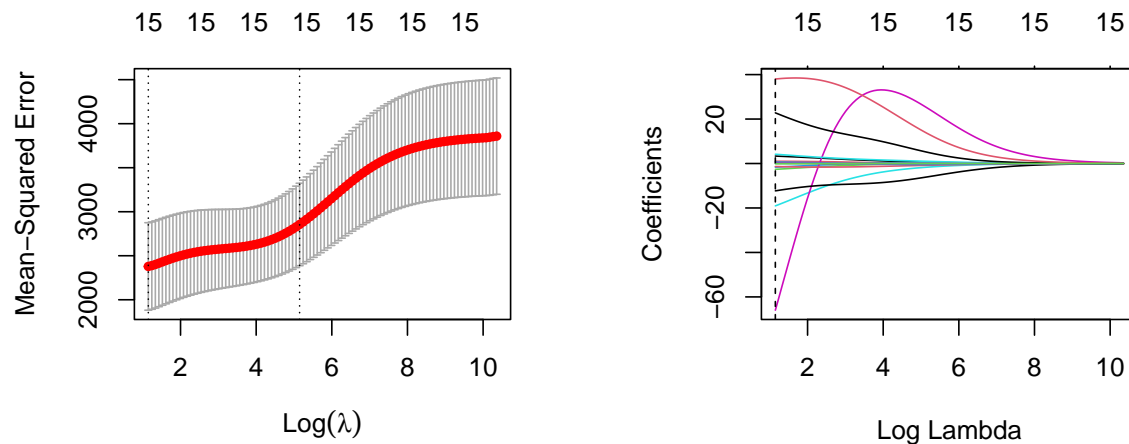
```
library(glmnet)
m.glm <- glm(mortality ~ .+HC:NOX, data=pollution)
X <- model.matrix(m.glm)[,-1]
m.ridge <- glmnet(x=X, y=pollution$mortality, alpha=0)
```

```
plot(m.ridge, xvar='lambda')
```



```
set.seed(222)
m.ridge.cv <- cv.glmnet(x=X, y=pollution$mortality, alpha=0)
```

```
par(mfrow=c(1,2))
plot(m.ridge.cv)
plot(m.ridge, xvar='lambda')
abline(v=log(m.ridge.cv$lambda.min), lty=2)
```




```
m.ridge.min <- glmnet(x=X, y=pollution$mortality, alpha=0,
                      lambda=m.ridge.cv$lambda.min)
cbind(coef(m.glm), coef(m.ridge.min))

## 16 x 2 sparse Matrix of class "dgCMatrix"
##                                     s0
## (Intercept)    1.901397e+03  1.463837e+03
## precipitation  4.160419e+00  3.355093e+00
## humidity       7.525489e-01  1.014058e+00
## Jan.temp       -3.944194e+00 -2.625912e+00
## July.temp      -2.363294e+00 -1.633865e+00
## over65         -2.722341e+01 -1.906860e+01
## house          -1.560127e+02 -6.612185e+01
## education      -1.583265e+01 -1.231362e+01
## comfort        -8.118738e-01 -1.357872e+00
## density        5.243647e-03  8.017681e-03
## office         2.298015e-01  6.291360e-01
## poor           5.801315e+00  4.177718e+00
## HC             7.736219e-01  1.324555e-01
## NOXSi          4.526420e+01  2.289480e+01
## SO2Si          3.148585e+01  3.798483e+01
## HC:NOXSi       -6.893134e-01 -7.057876e-02
```

Ridge regression confirms the association of NOX and SO2 on the mortality rate. We can use lasso for variable selection.

```
m.lasso <- glmnet(x=X, y=pollution$mortality, alpha=1)
```

```
plot(m.lasso, xvar='lambda')
```


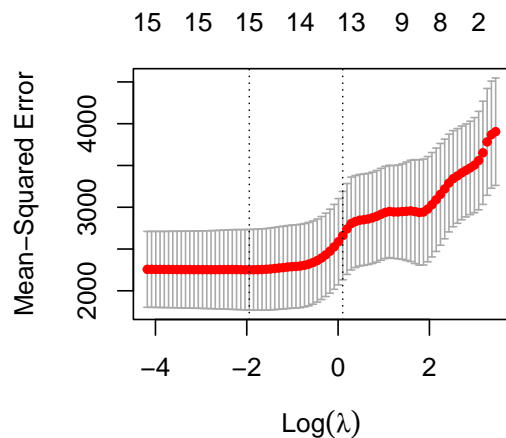


```
set.seed(222)
m.lasso.cv <- cv.glmnet(x=X, y=pollution$mortality, alpha=1)
m.lasso.cv

##
## Call:  cv.glmnet(x = X, y = pollution$mortality, alpha = 1)
##
## Measure: Mean-Squared Error
##
##      Lambda Index Measure      SE Nonzero
## min 0.1429    59    2252 482.3        15
## 1se 1.1067    37    2660 529.0        14
```

There is no relevant selection.

```
par(mfrow=c(1,2))
plot(m.lasso.cv)
plot(m.lasso, xvar='lambda')
abline(v=log(m.lasso.cv$lambda.min), lty=2)
```




```
m.lasso.min <- glmnet(x=X, y=pollution$mortality, alpha=1,
                      lambda=m.lasso.cv$lambda.min)
cbind(coef(m.glm), coef(m.ridge.min), coef(m.lasso.min))
```

```
## 16 x 3 sparse Matrix of class "dgCMatrix"
```

	s0	s0	s0
## (Intercept)	1.901397e+03	1.463837e+03	1.857920e+03
## precipitation	4.160419e+00	3.355093e+00	4.140763e+00
## humidity	7.525489e-01	1.014058e+00	8.675517e-01
## Jan.temp	-3.944194e+00	-2.625912e+00	-3.888083e+00
## July.temp	-2.363294e+00	-1.633865e+00	-2.393603e+00
## over65	-2.722341e+01	-1.906860e+01	-2.702569e+01
## house	-1.560127e+02	-6.612185e+01	-1.488006e+02
## education	-1.583265e+01	-1.231362e+01	-1.562861e+01
## comfort	-8.118738e-01	-1.357872e+00	-7.706245e-01
## density	5.243647e-03	8.017681e-03	5.882628e-03
## office	2.298015e-01	6.291360e-01	4.374290e-01
## poor	5.801315e+00	4.177718e+00	5.758559e+00
## HC	7.736219e-01	1.324555e-01	5.053675e-01
## NOXSi	4.526420e+01	2.289480e+01	3.846837e+01
## SO2Si	3.148585e+01	3.798483e+01	3.284600e+01
## HC:NOXSi	-6.893134e-01	-7.057876e-02	-4.187426e-01

```
par(mfrow=c(1,2))
plot(log(m.ridge$lambda), m.ridge$dev.ratio, type='l',
      xlab=expression(log(lambda)), ylab='Explained deviance -- ridge')
abline(v=log(m.ridge.cv$lambda.min), lty=2)
plot(log(m.lasso$lambda), m.lasso$dev.ratio, type='l',
      xlab=expression(log(lambda)), ylab='Explained deviance -- lasso')
abline(v=log(m.lasso.cv$lambda.min), lty=2)
```



Compare the results from ridge regression and lasso.

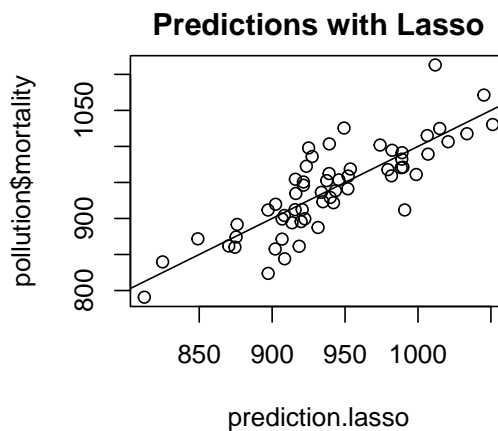

```
m.ridge$dev.ratio[m.ridge$lambda==m.ridge.cv$lambda.min]
## [1] 0.662835

m.lasso$dev.ratio[m.lasso$lambda==m.lasso.cv$lambda.min]
## [1] 0.691473
```

Lasso is preferable in terms of explained deviance.

```
prediction.ridge <- predict(m.ridge.min, newx=X)
prediction.lasso <- predict(m.lasso.min, newx=X)
```

```
par(mfrow=c(1,2))
plot(prediction.ridge, pollution$mortality,
      main='Predictions with Ridge')
abline(0,1)
plot(prediction.lasso, pollution$mortality,
      main='Predictions with Lasso')
abline(0,1)
```



```
min(m.ridge.cv$cvm)

## [1] 2378.412

min(m.lasso.cv$cvm)

## [1] 2251.617
```

We prefer lasso.

Compare the results with those from the linear model.

```
m.glm <- glm(mortality ~ .+HC:NOX, data=pollution)
library(boot)
set.seed(222)
m.glm.cv <- cv.glm(pollution, m.glm)
m.glm.cv$delta

## [1] 2516.973 2502.145
```

Lasso is preferable in terms of mean squared error. But since there is no substantial variable selection, it does not seem to be so interesting.

Try an automatic variable selection, using a forward procedure, adding the interactions suggested by the model of the first question.

```
library(leaps)
m.forward <- regsubsets(mortality ~ . + HC:NOX, data=pollution, nvmax=19,
  method='forward')
summary(m.forward)

## Subset selection object
## Call: regsubsets.formula(mortality ~ . + HC:NOX, data = pollution,
```

```

##      nvmax = 19, method = "forward")
## 15 Variables (and intercept)
##      Forced in Forced out
## precipitation      FALSE      FALSE
## humidity           FALSE      FALSE
## Jan.temp           FALSE      FALSE
## July.temp          FALSE      FALSE
## over65             FALSE      FALSE
## house              FALSE      FALSE
## education          FALSE      FALSE
## comfort            FALSE      FALSE
## density            FALSE      FALSE
## office             FALSE      FALSE
## poor               FALSE      FALSE
## HC                 FALSE      FALSE
## NOXSi              FALSE      FALSE
## SO2Si              FALSE      FALSE
## HC:NOXSi           FALSE      FALSE
## 1 subsets of each size up to 15
## Selection Algorithm: forward
##      precipitation humidity Jan.temp July.temp over65 house education comfort
## 1 ( 1 ) " " " " " " " " " " " "
## 2 ( 1 ) "*" " " " " " " " " " "
## 3 ( 1 ) "*" " " " " " " " " " "
## 4 ( 1 ) "*" " " " " " " "*" " " "
## 5 ( 1 ) "*" " " "*" " " "*" " " "
## 6 ( 1 ) "*" " " "*" " " "*" " " "
## 7 ( 1 ) "*" " " "*" " " "*" " " "
## 8 ( 1 ) "*" " " "*" "*" "*" " " "
## 9 ( 1 ) "*" " " "*" "*" "*" "*" " "
## 10 ( 1 ) "*" " " "*" "*" "*" "*" "*" "
## 11 ( 1 ) "*" " " "*" "*" "*" "*" "*" "
## 12 ( 1 ) "*" " " "*" "*" "*" "*" "*" "
## 13 ( 1 ) "*" "*" "*" "*" "*" "*" "*" "
## 14 ( 1 ) "*" "*" "*" "*" "*" "*" "*" "*"
## 15 ( 1 ) "*" "*" "*" "*" "*" "*" "*" "*"
##      density office poor HC NOXSi SO2Si HC:NOXSi
## 1 ( 1 ) " " " " " " " " " "
## 2 ( 1 ) " " " " " " " " " "
## 3 ( 1 ) " " " " " " " " "*" "
## 4 ( 1 ) " " " " " " " " "*" "
## 5 ( 1 ) " " " " " " " " "*" "
## 6 ( 1 ) "*" " " " " " " " " "*" "
## 7 ( 1 ) "*" " " "*" " " " " " "*" "

```



```
## 8 ( 1 ) "*" " " "*" " " " " "*" " "
## 9 ( 1 ) "*" " " "*" " " " " "*" " "
## 10 ( 1 ) "*" " " "*" " " "*" "*" " "
## 11 ( 1 ) "*" " " "*" "*" "*" "*" " "
## 12 ( 1 ) "*" " " "*" "*" "*" "*" "*"
## 13 ( 1 ) "*" " " "*" "*" "*" "*" "*"
## 14 ( 1 ) "*" " " "*" "*" "*" "*" "*"
## 15 ( 1 ) "*" "*" "*" "*" "*" "*" "*"

```

RSS criterion

```
summary(m.forward)$rss

## [1] 168671.67 148526.00 129820.13 115124.40 108617.31 102580.15 94046.62 85370.37
## [9] 79775.09 75398.52 74454.73 70875.16 70282.13 69978.05 69954.69

which.min(summary(m.forward)$rss)

## [1] 15

coef(m.forward, 15)

## (Intercept) precipitation humidity Jan.temp July.temp over65
## 1.901397e+03 4.160419e+00 7.525489e-01 -3.944194e+00 -2.363294e+00 -2.722341e+01
## house education comfort density office poor
## -1.560127e+02 -1.583265e+01 -8.118738e-01 5.243647e-03 2.298015e-01 5.801315e+00
## HC NOXSi SO2Si HC:NOXSi
## 7.736219e-01 4.526420e+01 3.148585e+01 -6.893134e-01

```

Adjusted R^2 criterion

```
which.max(summary(m.forward)$adjr2)

## [1] 12

coef(m.forward, 12)

## (Intercept) precipitation Jan.temp July.temp over65 house
## 1.941417e+03 4.182912e+00 -3.951313e+00 -3.006279e+00 -2.743225e+01 -1.595139e+02
## education density poor HC NOXSi SO2Si
## -1.595196e+01 4.670546e-03 6.650527e+00 8.299267e-01 4.663277e+01 3.012948e+01
## HC:NOXSi
## -7.676111e-01

```

BIC

```
which.min(summary(m.forward)$bic)
```


```
## [1] 8
```

```
coef(m.forward, 8)
```


```
## (Intercept) precipitation Jan.temp July.temp over65 education
## 1218.58525725 3.30033252 -2.27918048 -3.81679604 -17.45698605 -2.79524322
## density poor SO2Si
## 0.01093273 6.67409276 56.13858512
```

Graphical evaluation of the automatic selection

```
plot(m.forward)
```



```
plot(m.forward, scale='adjr2')
```



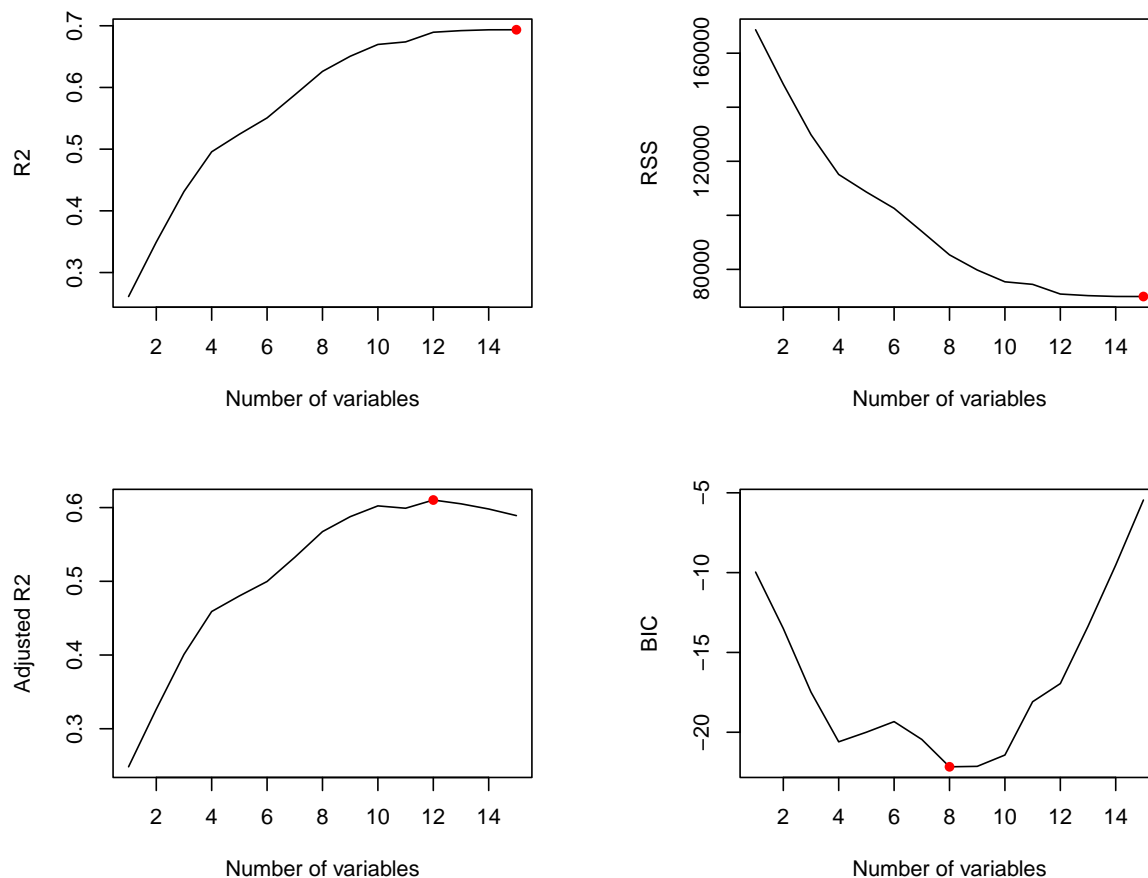
Selection based on BIC suggests an association between mortality with precipitation, poor, SO₂. Selection based on adjusted R^2 also adds density, HC and NOX.
Model ranking

```
par(mfrow=c(2,2))
## R2
plot(summary(m.forward)$rsq, xlab='Number of variables', ylab='R2', type='l')
## add the indication of the preferable model
points(which.max(summary(m.forward)$rsq),
summary(m.forward)$rsq[which.max(summary(m.forward)$rsq)], col='red', pch=16)
## RSS
plot(summary(m.forward)$rss, xlab='Number of variables', ylab='RSS', type='l' )
```

```

points(which.min(summary(m.forward)$rss),
summary(m.forward)$rss[which.min(summary(m.forward)$rss)], col='red', pch=16)
## adjusted R2
plot(summary(m.forward)$adjr2, xlab='Number of variables',
      ylab='Adjusted R2', type='l')
points(which.max(summary(m.forward)$adjr2),
summary(m.forward)$adjr2[which.max(summary(m.forward)$adjr2)],
      col='red', pch=16)
## BIC
plot(summary(m.forward)$bic, xlab='Number of variables', ylab='BIC', type='l')
points(which.min(summary(m.forward)$bic),
summary(m.forward)$bic[which.min(summary(m.forward)$bic)], col='red', pch=16)

```



Evaluate the regression model chosen by BIC.

```

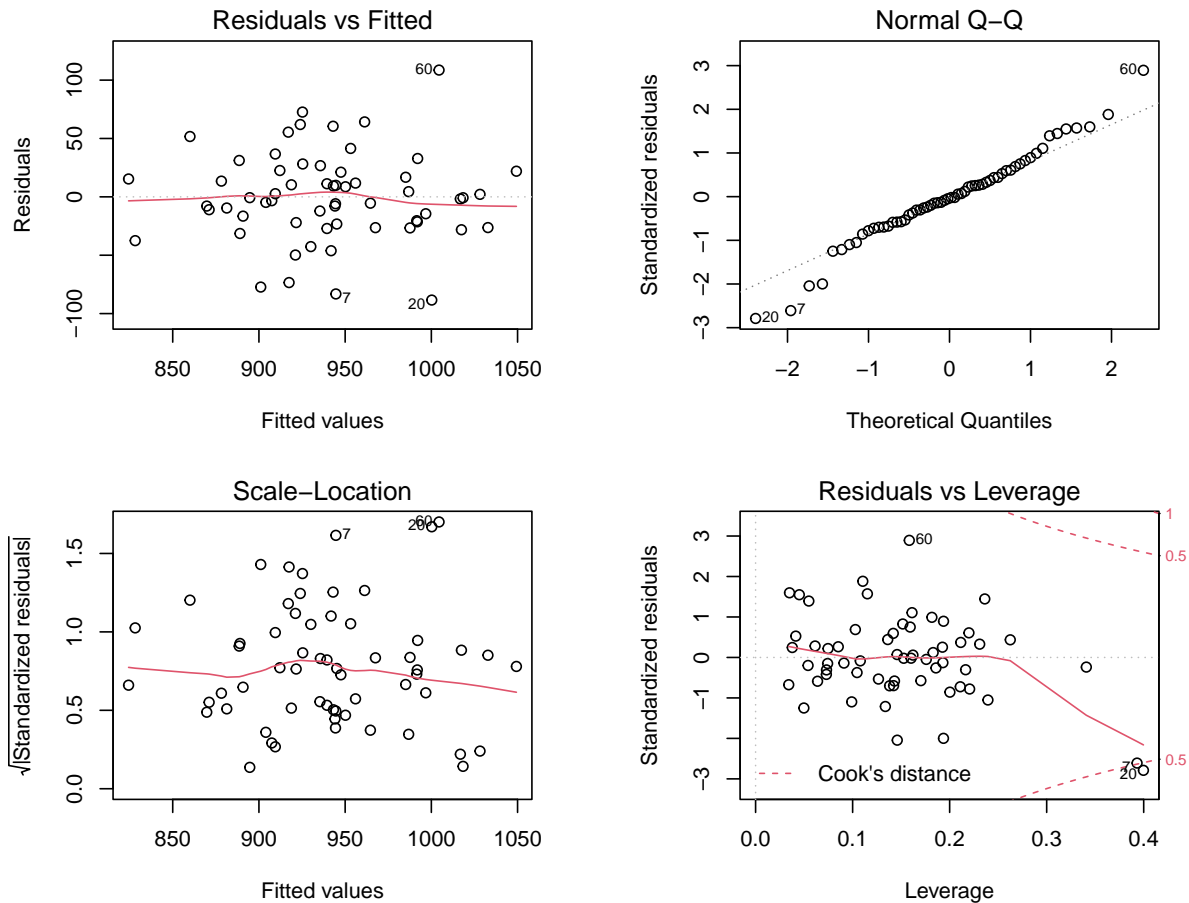
m.bic <- lm(mortality ~ precipitation + Jan.temp + July.temp +
            over65+education+density+poor +S02, data=pollution)
summary(m.bic)
##

```

```
## Call:
## lm(formula = mortality ~ precipitation + Jan.temp + July.temp +
##      over65 + education + density + poor + S02, data = pollution)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -88.411 -22.365  -1.283   21.336 108.597
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   1.219e+03  1.721e+02   7.082 4.03e-09 ***
## precipitation  3.300e+00  7.701e-01   4.286 8.09e-05 ***
## Jan.temp      -2.279e+00  7.475e-01  -3.049 0.003635 **
## July.temp     -3.817e+00  1.676e+00  -2.277 0.027035 *
## over65        -1.746e+01  4.665e+00  -3.742 0.000464 ***
## education     -2.795e+00  9.059e+00  -0.309 0.758910
## density        1.093e-02  4.137e-03   2.642 0.010903 *
## poor           6.674e+00  2.336e+00   2.857 0.006172 **
## S02Si          5.614e+01  1.808e+01   3.105 0.003107 **
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 40.91 on 51 degrees of freedom
## Multiple R-squared:  0.626, Adjusted R-squared:  0.5674
## F-statistic: 10.67 on 8 and 51 DF,  p-value: 1.165e-08
```

Check residuals

```
par(mfrow=c(2,2))
plot(m.bic)
```



We can check whether the principal component regression is useful.

```
library(pls)
set.seed(222)
m.pcr <- pcr(mortality ~ .+HC:NOX, scale=TRUE, validation='CV', data=pollution)
```

```
summary(m.pcr)

## Data:  X dimension: 60 15
## Y dimension: 60 1
## Fit method: svdpc
## Number of components considered: 15
##
## VALIDATION: RMSEP
## Cross-validated using 10 random segments.
##      (Intercept) 1 comps 2 comps 3 comps 4 comps 5 comps 6 comps 7 comps
## CV          62.73  58.26  55.34  52.39  52.72  53.69  52.80  52.27
## adjCV       62.73  58.02  54.94  52.07  52.45  53.41  52.34  51.94
##      8 comps 9 comps 10 comps 11 comps 12 comps 13 comps 14 comps 15 comps
```

```
## CV      53.91    51.89    58.56    58.97    63.07    69.82    50.54    49.53
## adjCV   53.57    51.99    57.43    57.80    61.76    68.22    49.63    48.63
##
## TRAINING: % variance explained
##          1 comps  2 comps  3 comps  4 comps  5 comps  6 comps  7 comps  8 comps
## X          29.66    46.78    60.78    69.94    78.22    84.3    88.64    92.08
## mortality  18.89    27.21    39.73    39.75    39.77    46.8    47.17    48.57
##          9 comps 10 comps 11 comps 12 comps 13 comps 14 comps 15 comps
## X          94.31    96.50    97.79    98.84    99.49    99.97    100.00
## mortality  50.31    55.99    56.53    56.93    57.17    68.34    69.36
```

How many principal components?

```
selectNcomp(m.pcr, method='onesigma', ncomp=15)


## [1] 2
```

How much variance is explained by the two first principal components?

```
explvar(m.pcr)

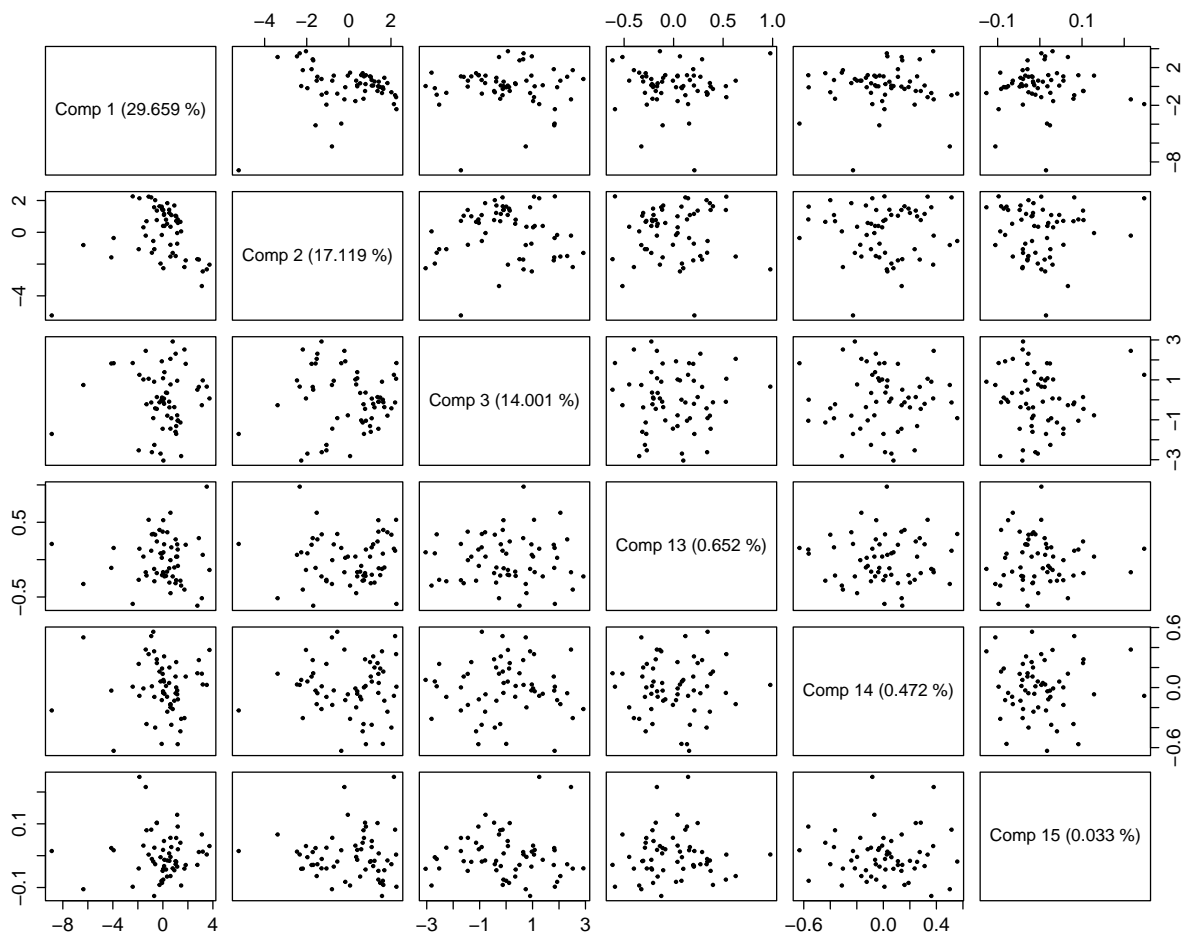
##      Comp 1      Comp 2      Comp 3      Comp 4      Comp 5      Comp 6      Comp 7      Comp
## 29.6592716 17.1187196 14.0014794  9.1564337  8.2889907  6.0750224  4.3419359  3.43718
##      Comp 9      Comp 10     Comp 11     Comp 12     Comp 13     Comp 14     Comp 15
##  2.2271121  2.1930932  1.2876528  1.0556113  0.6520201  0.4721439  0.0333317
```

```
coefplot(m.pcr, ncomp=1:5, legendpos='bottomleft', main='',
xlab='', ylab='regression coefficients', axes=FALSE)
axis(1, at=1:14, labels=colnames(pollution)[-1], las=2, cex=0.4, cex.axis=0.7)
axis(2, las=2, cex.axis=0.6)
```



On the basis of the first PC, mortality is associated to the temperature in July, to the features of the house (number of persons and comfort), to the poverty level of the area. The second PC gives weight also to the presence of pollutants NOX and SO2. Check for the presence of groups or outliers using the scoreplot about some of the principal components.

```
scoreplot(m.pcr, comps=c(1,2,3,13,14,15), cex=0.5, cex.lab=1.4, cex.axis=1.4, pch=19)
```

There is no evidence of problems.