

Regularization with R

Data Mining
Master Degree in Computer Science
University of Padova

a.y. 2017/2018

Annamaria Guolo

1 Hitters dataset

Consider dataset `Hitters` inside package `ISLR` we have already used when talking about model selection.

```
library(ISLR)
data(Hitters)
## omit missing data
hitters <- na.omit(Hitters)
dim(hitters)

## [1] 263 20

names(hitters)

## [1] "AtBat" "Hits" "HmRun" "Runs" "RBI" "Walks" "Years"
## [8] "CAtBat" "CHits" "CHmRun" "CRuns" "CRBI" "CWalks" "League"
## [15] "Division" "PutOuts" "Assists" "Errors" "Salary" "NewLeague"

hitters$Salary <- log(hitters$Salary)
## using the transformation applied in the previous analysis of the data
```

Consider a model relating the seasonal income to the other covariates. Ridge regression and lasso can be applied using functions inside library `glmnet`.

```
library(glmnet)
```

Start with ridge regression. We estimate the model through function `glmnet()`, which needs the matrix of covariates `X` and the vector of observations from `y`. Parameter `alpha` is set to 0 (indicator of ridge regression).

```
y <- hitters$Salary
X <- model.matrix(Salary ~ ., data=hitters)[,-1]
```

Function `model.matrix()` creates the matrix with covariates and in the meanwhile it transforms qualitative variables into dummies. Remember to eliminate the first column corresponding to the intercept.

```
m.ridge <- glmnet(X, y, alpha=0)
```

Output

```
m.ridge

##
## Call:  glmnet(x = X, y = y, alpha = 0)
##
##           Df        %Dev    Lambda
##  [1,] 19 7.070e-36 551.20000
##  [2,] 19 1.214e-02 502.20000
##  [3,] 19 1.330e-02 457.60000
##  [4,] 19 1.456e-02 417.00000
##  [5,] 19 1.595e-02 379.90000
##  [6,] 19 1.746e-02 346.20000
##  [7,] 19 1.911e-02 315.40000
##  [8,] 19 2.091e-02 287.40000
##  [9,] 19 2.288e-02 261.90000
## [10,] 19 2.503e-02 238.60000
## [11,] 19 2.736e-02 217.40000
## [12,] 19 2.990e-02 198.10000
## [13,] 19 3.267e-02 180.50000
## [14,] 19 3.568e-02 164.50000
## [15,] 19 3.895e-02 149.90000
## [16,] 19 4.249e-02 136.50000
## [17,] 19 4.634e-02 124.40000
## [18,] 19 5.050e-02 113.40000
## [19,] 19 5.500e-02 103.30000
## [20,] 19 5.985e-02  94.11000
## [21,] 19 6.509e-02  85.75000
## [22,] 19 7.073e-02  78.13000
## [23,] 19 7.679e-02  71.19000
## [24,] 19 8.329e-02  64.87000
## [25,] 19 9.025e-02  59.11000
## [26,] 19 9.768e-02  53.85000
## [27,] 19 1.056e-01  49.07000
## [28,] 19 1.140e-01  44.71000
## [29,] 19 1.229e-01  40.74000
## [30,] 19 1.324e-01  37.12000
```

```

## [31,] 19 1.423e-01 33.82000
## [32,] 19 1.528e-01 30.82000
## [33,] 19 1.637e-01 28.08000
## [34,] 19 1.751e-01 25.59000
## [35,] 19 1.869e-01 23.31000
## [36,] 19 1.991e-01 21.24000
## [37,] 19 2.117e-01 19.35000
## [38,] 19 2.247e-01 17.63000
## [39,] 19 2.378e-01 16.07000
## [40,] 19 2.512e-01 14.64000
## [41,] 19 2.647e-01 13.34000
## [42,] 19 2.783e-01 12.16000
## [43,] 19 2.918e-01 11.08000
## [44,] 19 3.053e-01 10.09000
## [45,] 19 3.186e-01 9.19500
## [46,] 19 3.317e-01 8.37800
## [47,] 19 3.444e-01 7.63400
## [48,] 19 3.568e-01 6.95600
## [49,] 19 3.687e-01 6.33800
## [50,] 19 3.802e-01 5.77500
## [51,] 19 3.912e-01 5.26200
## [52,] 19 4.016e-01 4.79400
## [53,] 19 4.114e-01 4.36800
## [54,] 19 4.206e-01 3.98000
## [55,] 19 4.293e-01 3.62700
## [56,] 19 4.374e-01 3.30400
## [57,] 19 4.448e-01 3.01100
## [58,] 19 4.518e-01 2.74300
## [59,] 19 4.582e-01 2.50000
## [60,] 19 4.641e-01 2.27800
## [61,] 19 4.695e-01 2.07500
## [62,] 19 4.745e-01 1.89100
## [63,] 19 4.790e-01 1.72300
## [64,] 19 4.832e-01 1.57000
## [65,] 19 4.871e-01 1.43000
## [66,] 19 4.906e-01 1.30300
## [67,] 19 4.939e-01 1.18800
## [68,] 19 4.969e-01 1.08200
## [69,] 19 4.997e-01 0.98590
## [70,] 19 5.023e-01 0.89830
## [71,] 19 5.048e-01 0.81850
## [72,] 19 5.070e-01 0.74580
## [73,] 19 5.091e-01 0.67960
## [74,] 19 5.111e-01 0.61920
## [75,] 19 5.130e-01 0.56420
## [76,] 19 5.148e-01 0.51410
## [77,] 19 5.165e-01 0.46840

```

```
## [78,] 19 5.181e-01 0.42680
## [79,] 19 5.197e-01 0.38890
## [80,] 19 5.212e-01 0.35430
## [81,] 19 5.227e-01 0.32280
## [82,] 19 5.241e-01 0.29420
## [83,] 19 5.255e-01 0.26800
## [84,] 19 5.268e-01 0.24420
## [85,] 19 5.281e-01 0.22250
## [86,] 19 5.294e-01 0.20280
## [87,] 19 5.307e-01 0.18470
## [88,] 19 5.319e-01 0.16830
## [89,] 19 5.331e-01 0.15340
## [90,] 19 5.343e-01 0.13980
## [91,] 19 5.355e-01 0.12730
## [92,] 19 5.366e-01 0.11600
## [93,] 19 5.377e-01 0.10570
## [94,] 19 5.388e-01 0.09633
## [95,] 19 5.399e-01 0.08777
## [96,] 19 5.409e-01 0.07997
## [97,] 19 5.420e-01 0.07287
## [98,] 19 5.430e-01 0.06639
## [99,] 19 5.439e-01 0.06050
## [100,] 19 5.449e-01 0.05512
```

The output reports the value for the deviance for each value of λ . Actually, the estimated object includes many other quantities

```
names(m.ridge)

## [1] "a0"      "beta"    "df"      "dim"     "lambda"  "dev.ratio" "nulldev"
## [8] "npasses" "jerr"    "offset"  "call"    "nobs"
```

- `a0`: estimated intercept for each model fitted with a different λ
- `beta`: $p \times (\text{number of } \lambda)$ matrix with the estimates of the coefficients
- `lambda`: values of λ
- `dev.ratio`: 1- model deviance/null deviance;
- `nulldev`: null deviance


How many λ are considered?

```
length(m.ridge$lambda)

## [1] 100
```


Graphical evaluation of the coefficients associated to the covariates

```
plot(m.ridge, xvar='lambda', xlab=expression(log(lambda)))
```



Option `xvar='lambda'` specifies that the x-axis is expressed in terms of λ . Alternatives are deviance values and L1-norm values. See, for example

```
plot(m.ridge, xlab=expression(log(lambda)))
```



Option `xlab=expression(log(lambda))` insert the mathematical symbol for λ in the axis. Numbers (19, repeated) over the graph indicate the number of covariates entering the model as λ varies: 19 is repeated, as ridge regression is not a selection method.

Look for the best λ using cross validation, using function `cv.glmnet()`, with a syntax similar to that in `glmnet()`.

Fix the seed

```
set.seed(2906)
```


```
cv.ridge <- cv.glmnet(X, y, alpha=0)
```

For default, R considers 10-fold cross validation. The resulting object includes several quantities, such as

- `lambda`: values of λ
- `cvm`: the MSE for each λ
- `cvsd`: the estimate of the standard error of `cvm`
- `lambda.min`: the value of λ associated to the minimum `cvm`
- `lambda.1se`: the values of λ associated to the minimum `cvm` within 1 standard error.

Graphical representation

```
plot(cv.ridge, xlab=expression(log(lambda)))
```



The plots shows the values of `cvm` for each $\log(\lambda)$ together with the associated confidence interval. The two dashed lines are the values of `log-lambda.min` and `log-lambda.1se`.

λ from cross validation

```
best.lambda <- cv.ridge$lambda.min
best.lambda
## [1] 0.060496
```

Find the minimum MSE

```

cv.ridge$cvm[cv.ridge$lambda==best.lambda]

## [1] 0.403175

## or, equivalently,
min(cv.ridge$cvm)

## [1] 0.403175

```

Re-estimate the model using the best λ

```

m.ridge.min <- glmnet(X, y, alpha=0, lambda=best.lambda)
m.ridge.min

##
## Call:  glmnet(x = X, y = y, alpha = 0, lambda = best.lambda)
##
##          Df    %Dev Lambda
## [1,] 19 0.5439 0.0605

```

Coefficients of the model

```

coef(m.ridge.min)

## 20 x 1 sparse Matrix of class "dgCMatrix"
##                               s0
## (Intercept)  4.595305e+00
## AtBat        -5.160388e-04
## Hits         4.396714e-03
## HmRun        3.359157e-03
## Runs         3.110850e-03
## RBI          7.308564e-04
## Walks        5.371948e-03
## Years        4.171532e-02
## CAtBat       3.892745e-05
## CHits        1.913195e-04
## CHmRun       -5.602232e-05
## CRuns        2.972823e-04
## CRBI         1.652623e-04
## CWalks       -3.891595e-04
## LeagueN      2.052706e-01
## DivisionW    -1.666745e-01
## PutOuts      2.925873e-04
## Assists      3.949817e-04
## Errors       -1.161143e-02
## NewLeagueN  -1.060981e-01


```

Graphical representation of the coefficients for the best λ and model deviance

```

par(mfrow=c(1,2))
plot(m.ridge, xvar='lambda', xlab=expression(log(lambda)))
## add on the line corresponding to the best lambda
abline(v=log(best.lambda), lty=2)
## deviance
plot(log(m.ridge$lambda), m.ridge$dev.ratio, type='l',
      xlab=expression(log(lambda)), ylab='Explained deviance')
abline(v=log(best.lambda), lty=2)

```



The maximum explained deviance is obtained for the minimum (best) λ and it is equal to

```

max(m.ridge$dev.ratio)
## [1] 0.5448529


```

Now move to lasso. The syntax is similar to that for ridge regression, but specifying $\alpha=1$

```
m.lasso <- glmnet(X, y, alpha=1)
```

Graphical representation of the coefficients

```
plot(m.lasso, xvar='lambda', xlab=expression(log(lambda)))
```

Look for λ that minimizes the MSE

```
## fix the seed to the same value used for ridge regression
set.seed(2906)
cv.lasso <- cv.glmnet(X, y, alpha=1)
```

Minimum λ from cross validation

```
best.lambda.lasso <- cv.lasso$lambda.min
```

Minimum MSE

```
min(cv.lasso$cvm)
## [1] 0.4026953
```

On the basis of MSE, the model fitted with lasso is preferable. In addition, the resulting model with lasso is simplest.

Re-estimate the model using the best λ from cross-validation

```
m.lasso.min <- glmnet(X, y, alpha=1, lambda=best.lambda.lasso)
```

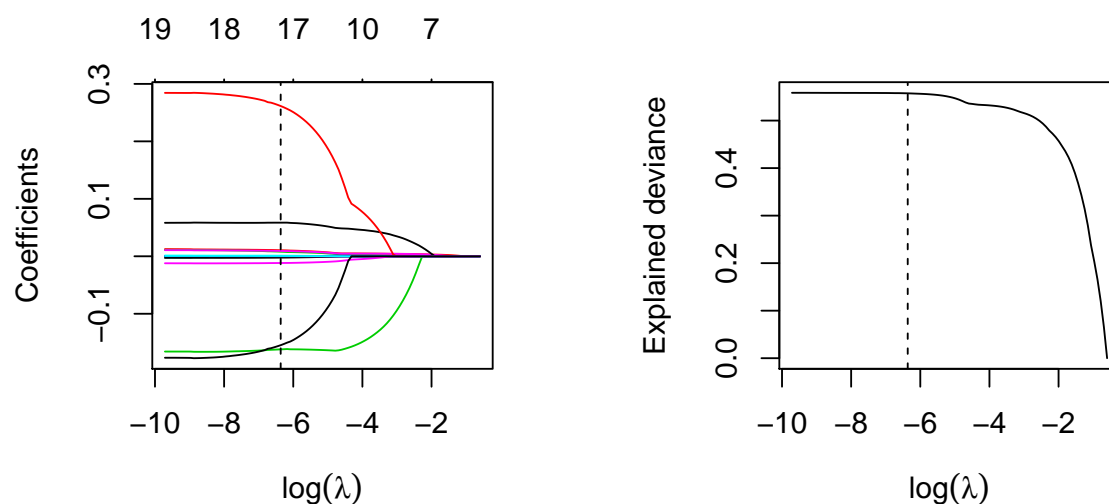
Coefficients

```
coef(m.lasso.min)
## 20 x 1 sparse Matrix of class "dgCMatrix"
##              s0
## (Intercept)  4.593462e+00
## AtBat       -2.311993e-03
## Hits        1.056073e-02
```

```
## HmRun      7.973244e-03
## Runs       .
## RBI        -2.285376e-04
## Walks      9.294713e-03
## Years      5.881411e-02
## CAtBat     2.017179e-05
## CHits      4.843161e-05
## CHmRun     2.194080e-04
## CRuns      1.068970e-03
## CRBI       .
## CWalks     -1.089961e-03
## LeagueN    2.615124e-01
## DivisionW  -1.616904e-01
## PutOuts    3.204166e-04
## Assists    5.616621e-04
## Errors     -1.166310e-02
## NewLeagueN -1.543706e-01
```

Some of the coefficients are zero, so the lasso performed a model selection.
Graphical representation of the coefficients for the best λ and model deviance

```
par(mfrow=c(1,2))
plot(m.lasso, xvar='lambda', xlab=expression(log(lambda)))
## add on the line corresponding to the best lambda
abline(v=log(best.lambda.lasso), lty=2)
## deviance
plot(log(m.lasso$lambda), m.lasso$dev.ratio, type='l',
      xlab=expression(log(lambda)), ylab='Explained deviance')
abline(v=log(best.lambda.lasso), lty=2)
```



The maximum explained deviance is obtained for the minimum (best) λ and it is equal to

```
max(m.lasso$dev.ratio)

## [1] 0.5584737
```

2 Leukemia dataset

Consider the Leukemia data about the gene expression in cancer cells obtained from 72 subjects with acute myeloid leukemia and acute lymphoblastic leukemia. Data are available in the R workspace *Leukemia.RData*. Upload the data

```
load("Leukemia.RData")
ls()

## [1] "best.lambda"      "best.lambda.lasso" "cv.lasso"          "cv.ridge"
## [5] "hitters"          "Hitters"           "Leukemia"          "m.lasso"
## [9] "m.lasso.min"      "m.ridge"           "m.ridge.min"       "X"
## [13] "y"

names(Leukemia)

## [1] "x" "y"
```

There are the objects *x* and *y*. Response *y* is categorical (diseased/nondiseased)

```
table(Leukemia$y)

##
##  0  1
## 47 25
```

Object *x* is the matrix with the observations for the covariates

```
dim(Leukemia$x)

## [1] 72 3571
```

There are 3571 covariates. Clearly, a standard logistic regression model cannot work here. In fact, look at the output


```
m <- glm(y ~ x, data=Leukemia, family='binomial')
```

here not reported for space reason. Use ridge regression and lasso instead. Start with ridge regression

```
leukemia.ridge <- glmnet(Leukemia$x, Leukemia$y, alpha=0, family='binomial')
```

Note that we specify family='binomial' as y is a binary indicator.

```
plot(leukemia.ridge, xvar='lambda', xlab=expression(log(lambda)))
```



Which values of λ are used?

```
summary(leukemia.ridge$lambda)
```

```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##  4.093  12.946  40.942  89.195 129.462 409.310
```

Select the best value of λ using cross validation, previously extending the grid of values of λ through option lambda.min


```
set.seed(111)
cv.leukemia.ridge <- cv.glmnet(Leukemia$x, Leukemia$y, alpha=0, family='binomial',
                               lambda.min = 1e-4)
best.lambda.leukemia <- cv.leukemia.ridge$lambda.min
best.lambda.leukemia
```

```
## [1] 0.0449217
```

```
min(cv.leukemia.ridge$cvm)
```

```
## [1] 0.1708827
```

```
plot(cv.leukemia.ridge)
```




Re-estimate the model using the best λ chosen from cross-validation

```
leukemia.ridge.min <- glmnet(Leukemia$x, Leukemia$y, alpha=0, family='binomial',
                             lambda=best.lambda.leukemia)
```

Graphical representation of the coefficients for the best λ and model deviance. Remember to extend the grid of values of λ

```
leukemia.ridge <- glmnet(Leukemia$x, Leukemia$y, alpha=0, family='binomial',
                         lambda.min = 1e-4)
```

```
par(mfrow=c(1,2))
plot(leukemia.ridge, xvar='lambda', xlab=expression(log(lambda)))
## add on the line corresponding to the best lambda
abline(v=log(best.lambda.leukemia), lty=2)
## deviance
plot(log(leukemia.ridge$lambda), leukemia.ridge$dev.ratio, type='l',
      xlab=expression(log(lambda)), ylab='Explained deviance')
abline(v=log(best.lambda.leukemia), lty=2)
```



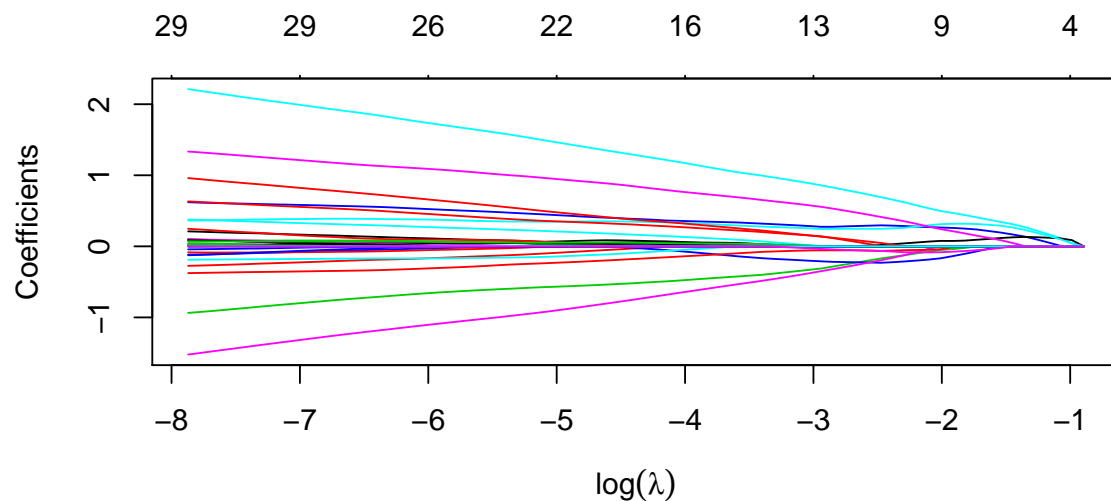
The maximum explained deviance is obtained for the minimum (best) λ and it is equal to

```
max(leukemia.ridge$dev.ratio)
## [1] 0.9975967
```

Lasso

```
leukemia.lasso <- glmnet(Leukemia$x, Leukemia$y, alpha=1, family='binomial',
  lambda.min = 1e-4)
```

```
plot(leukemia.lasso, xvar='lambda', xlab=expression(log(lambda)))
```



Select λ from cross validation

```

set.seed(111)
cv.leukemia.lasso <- cv.glmnet(Leukemia$x, Leukemia$y, alpha=1, family='binomial',
                               lambda.min = 1e-4)
best.lambda.leukemia.lasso <- cv.leukemia.lasso$lambda.min
best.lambda.leukemia.lasso


## [1] 0.0005538157

min(cv.leukemia.lasso$cvm)

## [1] 0.1443071

```

```
plot(cv.leukemia.lasso)
```



Re-estimate the model using the best λ from cross-validation

```

leukemia.lasso.min <- glmnet(Leukemia$x, Leukemia$y, alpha=1, family='binomial',
                             lambda=best.lambda.leukemia.lasso)

```

Pay attention to the different time consumption...

How many coefficients are set equal to zero? None in ridge regression.

```

id.zero <- which(coef(leukemia.lasso.min)==0)
length(id.zero)

## [1] 3541

nonzero <- length(coef(leukemia.lasso.min))-length(id.zero)
nonzero

## [1] 31

```

There are 3541 zero coefficients, so lasso selects only 30 variables (we need to eliminate the intercept). The chosen variables are

```
id.nonzero <- which(coef(leukemia.lasso.min)!=0)
varnames <- rownames(coef(leukemia.lasso.min))[id.nonzero]
values <- coef(leukemia.lasso.min)[id.nonzero]
names(values) <- varnames
values
```

##	(Intercept)	V158	V219	V456	V657	V672
##	-2.8824214763	0.0018961082	-0.2285781638	-0.9231795133	-0.1237991380	-1.4074641347
##	V888	V918	V926	V956	V979	V1007
##	0.1704091156	0.2085950745	0.0317573289	0.6160689383	2.0840545978	0.0354589050
##	V1219	V1569	V1652	V1796	V1835	V1946
##	-0.3447530778	0.0101644338	0.3413735697	0.0001313399	0.0646048792	0.9194201601
##	V2230	V2239	V2481	V2727	V2831	V2859
##	0.1121487684	-0.1144706536	1.2448989053	-0.1306291781	0.0066234777	-0.0629515620
##	V2888	V2929	V3038	V3098	V3125	V3158
##	0.3532643009	0.0902908966	0.1319079465	0.5965100779	0.0142853358	0.0783746529
##	V3181					
##	-0.1169728169					

Try to see what happens when using λ equal to `lambda.1se` in place of `lambda.min` from cross-validation.